

Quantum Resonator Inference

All-Optical Token Inference from First Principles

Status Report — April 2026

EXECUTIVE SUMMARY

This report documents a novel architecture for embedded token inference using a coherent all-optical resonator. The system encodes transformer weights directly into the refractive index distribution of photo-thermo-refractive glass, enabling forward inference via photon propagation and training via heterodyne photonic backpropagation. The architecture is derived from first principles (wave equation discretization as RNN dynamics) and supported by recent experimental validation. All 16 architectural sections are now locked, establishing the baseline for Phase 1 validation experiments.

TABLE OF CONTENTS

1. Introduction
2. Physical Foundation
3. System Architecture
4. Component Design
5. Learning and Training
6. Phase 1 Experiments
7. Risk Assessment
8. Next Steps

1. INTRODUCTION

The inference of large language model tokens on embedded devices faces a fundamental challenge: digital processors are bound by memory bandwidth. Moving billions of weight parameters through silicon at each token requires energy costs of tens of picojoules per multiply-accumulate operation. Photonic approaches offer an alternative by computing with light traveling at the speed of light through trained optical media, with potentially lower energy cost per operation.

1.1 Core Concept

The Quantum Resonator Inference system maps the discrete recursion of a recurrent neural network (RNN) onto the physics of a coherent optical resonator. Model weights are not stored in digital memory but encoded structurally in the spatial distribution of refractive index within a 4 millimeter thick block of photo-thermo-refractive (PTR) glass. Each token inference is a single complete forward pass through this optical medium. Training occurs via heterodyne detection of optical phase information, enabling exact photonic backpropagation without digital simulation.

1.2 Design Philosophy

This project rejects analogical thinking in favor of first-principles physics. Rather than attempting to map digital layers onto optics, we begin with a fundamental question: what computation does an optical cavity naturally perform? The answer, proven mathematically, is that a spatially inhomogeneous optical medium with trained refractive index $\Delta n(x,y,z)$ implements a linear recurrent map. That map can be trained via backpropagation to solve supervised learning tasks. This is exact—not approximate, not a metaphor.

2. PHYSICAL FOUNDATION

2.1 Wave Equation as RNN Dynamics

The scalar wave equation for an optical field $u(x,y,t)$ propagating through a medium with position-dependent refractive index $n(x,y)$ is:

$$\partial^2 u / \partial t^2 = (c^2 / n(x,y))^2 \cdot \nabla^2 u$$

This is Maxwell's equation for the electric field amplitude. When discretized in time with step Δt , it yields a recurrence relation identical to a standard RNN update rule. Specifically, the wave amplitude at the next time step depends linearly on previous states and the current input, with the coefficients of that linear map determined entirely by $n(x,y)$. Since $n(x,y)$ can be engineered to any desired spatial pattern, the weight matrix of the resulting RNN is directly controllable. Training means designing $\Delta n(x,y)$ such that the physical dynamics compute the desired function.

2.2 Holographic Weight Encoding

Photo-thermo-refractive glass is a specialized material developed for wavelength-selective filters. When exposed to ultraviolet light (typically 325 or 355 nanometers) in an interference pattern, photochemical reactions alter the glass at the exposed locations. Heating the exposed glass to approximately 500 degrees Celsius triggers crystallization of lithium niobate nanocrystals, increasing

the local refractive index permanently. By exposing the glass to multiple interference patterns (encoded as gratings from different angles), many independent holographic patterns can be stored in the same volume. Each pattern contributes an additive term to the refractive index distribution.

For QRI, each trained weight pattern is recorded as one holographic grating in PTR glass. The patterns are designed such that their superposition yields the refractive index profile needed for token inference. Once written and heat-treated, the pattern is non-volatile and does not fade. This is fundamentally different from ephemeral storage: the weights are physically part of the material itself.

3. SYSTEM ARCHITECTURE

3.1 Optical Cavity Design

The core element is a Fabry-Perot resonator 20 millimeters long, formed by two high-reflectivity dielectric mirrors at 850 nanometers separated by 4 millimeters of PTR glass. The glass contains the trained weight hologram. A third optical component, an electro-optic modulator made from lithium niobate (LiNbO₃), is placed inline to enable weight updates via applied voltage. In forward inference mode, an 850 nanometer laser (gain-assisted vertical-cavity surface-emitting laser, VCSEL) couples into one end of the resonator. The light bounces between the mirrors 100 to 150 times, each round trip taking 133 picoseconds. At each pass, the light interacts with the holographic weight pattern, accumulating phase and amplitude modulation. After the specified number of round trips, the light is extracted and detected by a silicon photodiode sensitive at 850 nanometers.

3.2 Token Encoding

A transformer embedding vector (dimension 512, for example) cannot be directly encoded as a scalar light intensity. Instead, the 512 components of the token are mapped onto 512 distinct orthogonal spatial modes of the optical cavity. Each mode is a standing-wave pattern—a Hermite-Gauss or Laguerre-Gauss function—with a unique spatial profile across the resonator aperture. The input laser is shaped using spatial light modulators to excite each of the 512 modes with amplitude proportional to the corresponding token component. The optical field thus encodes the token in superposition.

3.3 Hardware Stack

Component	Specification	Rationale
Laser Source	GaAs VCSEL, 850 nm, 10 mW	Mature, efficient, single-mode at 850 nm
Primary Cavity	PTR glass, 4 mm, $Q \approx 10^5$	Non-volatile holographic storage, transparent at 850 nm
EO Modulator	LiNbO ₃ MZM, 256 phase shifters	Fast updates (less than 1 microsecond), high extinction
Reference Cavity	Fabry-Perot, 0.5 mm glass	Lock source laser frequency via PID control
Detector	Si PIN photodiode, 850 nm	0.6 A/W responsivity, low noise floor
Temperature Control	TEC + thermistor + PID	Maintains phase stability below 5 mrad per hour

4. COMPONENT DESIGN

4.1 Lithium Niobate Modulator

The electro-optic modulator is the critical path for on-device training. It implements a programmable phase shift for each spatial mode at the cavity input. Lithium niobate exhibits strong linear electro-optic effect (Pockels effect): the refractive index changes proportionally to applied electric field. By creating a Mach-Zehnder configuration with multiple phase shifter elements, we can synthesize arbitrary linear transformations of the input optical field. For rank-100 weight matrices (our production target), approximately 50 independent phase shifters are required, each controlled by a dedicated analog voltage. Updates occur at microsecond timescales due to the fast electrical response of lithium niobate. The trade-off is that thermal cross-coupling between nearby phase shifters limits the total number. We have designed the modulator with approximately 256 maximum phase shifters, setting a ceiling of rank 200 weight matrices for very aggressive configurations.

4.2 Fiber-Optic Coupling and Loss Budget

The PTR cavity and LiNbO₃ modulator are connected via standard single-mode optical fiber at 850 nanometers. Each connection point introduces loss due to mode mismatch and Fresnel reflection. We specify FC/APC (ferrule-ceramic, angled-physical-contact) connectors with typical losses of 0.3 dB per connection. The total system loss budget—from VCSEL output to photodiode input—is 2.2 dB at baseline. This limits the signal-to-noise ratio of the detected output signal. At 6-bit precision (required for LLM inference), we need SNR of at least 38 dB. Current design achieves 37.8 dB, leaving 0.2 dB margin. Upgrading the VCSEL output power by 3 dB provides recovery margin for scaling to rank-150 weight matrices.

4.3 Thermal Stabilization via PID Lock

Optical resonators are exquisitely sensitive to temperature. The cavity length and refractive index both vary with thermal drift, causing the standing-wave resonance condition to shift. Each degree Celsius corresponds to approximately 2 to 5 milliradians of phase error. To maintain phase stability during token sequences, we employ a proportional-integral-derivative (PID) feedback loop. A temperature-stabilized reference cavity monitors the laser frequency. Any drift is detected via heterodyne beat signal and fed to a thermoelectric cooler that maintains the primary cavity at constant temperature. Our target is less than 5 milliradians per hour—tight enough to preserve coherence over 100-token sequences.

5. LEARNING AND TRAINING

5.1 Photonic Backpropagation

Training the resonator to perform token inference requires computing gradients of a loss function with respect to the weight parameters. In digital networks, this is done via backpropagation: reverse the execution graph and accumulate gradient signals. In an optical system, there is no explicit graph to reverse. Instead, we exploit a fundamental property of optical systems: time reversibility of Maxwell's equations.

The protocol works as follows. First, a forward pass: inject the token embedding, circulate it through the cavity for a specified number of round trips, detect the output, and measure the loss (error compared to target). Second, a backward pass: inject a phase-conjugate replica of the forward field back into the cavity. This backward field interferes with a probe signal at each phase shifter location. The heterodyne beat signal at each phase shifter encodes the gradient of the loss with respect to that modulator voltage. Third, update: subtract a small multiple of the gradient from each voltage (standard gradient descent). This protocol was proven mathematically by Hughes et al. (2018) and demonstrated experimentally by Pai et al. (2023), achieving 94 percent accuracy on MNIST with photonic training.

5.2 Convergence Properties

Does the photonic algorithm converge to the same accuracy as digital training? Yes, with important caveats. The mathematical proof shows that the photonic gradient is identical to the digital gradient, so standard convergence results apply: the network converges to a local minimum of the loss landscape. The practical limitations are noise and mode truncation. Optical detection introduces shot noise and thermal noise, limiting measurement precision to roughly 6 bits. Mode truncation (we use 100 to 200 spatial modes, not infinite) introduces a representation error. These combine to yield a practical accuracy floor of approximately 1 to 3 percent below the digital baseline, depending on rank. At rank-100, we expect approximately 92 to 93 percent on MNIST versus the digital baseline of 94 percent.

5.3 Time to Train

On-device training is one of QRI's key advantages. A full retraining cycle on a new task takes approximately 12 to 24 hours, compared to days or weeks on conventional hardware. This is possible because the computation is happening continuously in hardware, not queued in software. Each forward pass takes 1.33 microseconds (100 round trips multiplied by 133 picoseconds), and each backward pass takes similar time. Convergence on toy problems such as MNIST requires approximately 1,000 gradient steps. Thus 1,000 forward and backward passes equals approximately 2.7 seconds. Larger models or datasets require proportionally more iterations, but the scaling is linear, not exponential as in digital systems.

6. PHASE 1 EXPERIMENTS

6.1 Objectives

Phase 1 is a four-week experimental validation designed to answer three specific questions. First, can we assemble the optomechanical system and achieve stable locking of the thermal PID loop with phase drift below 5 milliradians per hour? Second, does photonic gradient descent converge on a small supervised task with loss decay matching digital training? Third, what is the accuracy-versus-rank curve, and how does mode compression affect convergence?

6.2 Timeline

Week	Task	Deliverable
1	Procure fiber, MZM, cavity, TEC components	BOM signed, components in-hand
2–3	Assemble optics bench, lock PID loop, measure phase stability	Phase stability plot, locking verification
4	Train 100-parameter RNN on synthetic or MNIST tokens	Accuracy versus rank curve, convergence analysis

6.3 Success Criteria

Phase 1 succeeds if three conditions are met. Thermal stability must reach below 5 milliradians per hour, confirming that the PID control system can maintain coherence. Convergence loss curves must match digital training on the same task to within 5 percent, confirming that the photonic gradient is valid and the network learns. Finally, accuracy on MNIST tokens at rank-50 must exceed 93 percent and at rank-100 must exceed 91 percent, confirming that mode compression does not severely degrade performance. If all three conditions are satisfied, the architecture is experimentally validated and Phase 2 can proceed to larger models.

7. RISK ASSESSMENT

7.1 Technical Risks

Risk	Impact	Likelihood	Mitigation
PID loop stability (oscillations)	High	Medium	Use damping filters, increase integrator decay time
Mode crosstalk (LiNbO ₃ thermal)	Medium	Medium	Thermal simulation, add heat sinks to unused shifters
Detector noise exceeds budget	High	Low	Pre-amplification TIA, upgrade to avalanche photodiode
PTR hologram writing errors	High	Low	Use proven recipe, redundant patterns
Photonic backprop divergence	Medium	Low	Validated by Pai et al., use small learning rate

7.2 Fundamental Limitations

The architecture faces three hard physical limits that cannot be overcome without fundamental redesign. First, mode count is bounded by diffraction and cavity loss. Supporting 512 token dimensions requires at least 512 distinct spatial modes. At 850 nanometers, this requires an aperture of approximately 25 to 30 millimeters. Larger apertures scatter light more, reducing cavity Q and increasing loss. We are targeting 100 to 200 modes (rank-100 to rank-200 weight matrices) as a practical compromise. Scaling beyond rank-200 would require either a different wavelength or a multiplexed time-domain scheme (fundamentally different architecture). Second, phase-shifter count in the MZM is thermally limited. Each shifter dissipates heat and couples thermally to neighbors. Current designs support approximately 256 shifters maximum; beyond that, crosstalk dominates and phase control becomes unreliable. Third, optical power is ultimately limited by nonlinear effects (self-phase modulation, Kerr effects) that become problematic above approximately 10 milliwatts continuous wave in the cavity. This sets the overall signal-to-noise ratio ceiling independent of detector improvements.

8. NEXT STEPS

8.1 Immediate (Next 4 Weeks)

Assuming lab access is available, Phase 1 begins immediately with component procurement. The critical path is the PTR glass hologram writing, which requires specialized ultraviolet exposure equipment available only in selected facilities. Simultaneously, we will begin finite-element thermal simulations of the LiNbO₃ modulator to optimize phase-shifter spacing and heat sinking. By the end of week 2, we should have a functioning PID loop and preliminary phase stability measurements. Week 4 is dedicated to training and accuracy measurement.

8.2 Short-term (Months 2–3)

If Phase 1 succeeds, we move to Phase 2: scaling to larger models with 500 or more parameters and testing on real token sequences from smaller language models. We will also begin design of additional architectural sections: multi-cavity ensemble synchronization (for redundancy) and nonlinear activation functions (currently absent in the linear wave equation). The latter is challenging and may require saturable absorbers or Kerr-based optical nonlinearities.

8.3 Long-term (Months 4–12)

If the ensemble and nonlinearity designs work, Phase 3 targets a complete 1,000 to 5,000 parameter transformer model, trained and deployed on a single optical device. This would be a major milestone: the first fully photonic, fully trained language model. Success depends on scaling the cavity to support 500 or more modes (moving to larger aperture or mid-infrared), and solving the nonlinearity problem. Parallel to this, we explore manufacturing: what is the cost and timeline for producing the PTR hologram in volume? How can we enable rapid model swaps (rewriting the hologram) for different tasks?