

Quantum Resonator Inference

System Design Document

Jason Edelman

2026-04-28

Quantum Resonator Inference: System Design Document

Quantum Resonator Inference — Architecture Specification

1. Problem Statement
 2. Design Principles
 3. Theoretical Foundation (ARCH-1)
 - 3.1 Wave Equation as RNN (Hughes et al. 2019)
 - 3.2 Holographic Weight Storage (Psaltis et al. 1990)
 - 3.3 From Scalar Wave to Token Embedding
 4. System Architecture
 - 4.1 System Overview
 - 4.2 Resonator Geometry (ARCH-2, ARCH-3)
 - 4.3 The Pure-Glass Resonator
 - 4.4 Input Encoding (ARCH-3)
 - 4.5 Inter-Layer Coupling (ARCH-8)
 - 4.6 Activation Function (ARCH-9, locked 2026-04-27)
 - 4.7 SNR and Noise Budget (ARCH-5)
 - 4.8 Token Throughput (ARCH-4)
 - 4.9 Holographic Weight Capacity (ARCH-7)
 - 4.10 Thermal Management (ARCH-10)
 5. Training Architecture
 - 5.1 The Weight Translation Problem (ARCH-11)
 - 5.2 Wavelength Separation: Write at 532nm, Infer at 850nm (ARCH-11)
 - 5.3 In-Situ Training Protocol (ARCH-11)
 - 5.4 Grating Stability During Inference (ARCH-11)
 - 5.5 Batch Size and Epoch Scale (ARCH-11)
 6. Scaling Architecture
 - 6.1 Block Duplication (ARCH-17)
 - 6.2 Clone-and-Fine-Tune for Scale Training (ARCH-17)
 - 6.3 Deployment Model (ARCH-12, ARCH-13)
 7. Rank Scaling (ARCH-16)
 - 7.1 Rank and Mode Basis
 - 7.2 Rank Scaling Regimes
 - 7.3 Rank-Loss Tradeoff
 8. Convergence Theory (ARCH-14, ARCH-15)
 - 8.1 Photonic Backpropagation (ARCH-14)
 - 8.2 Loss Landscape (ARCH-15)
 9. Decision Log
 10. Open Experiments
 - 10.1 EXP-7: In-Situ Training Convergence and Clone Validation (Expanded)
 11. Architecture Summary
- Architecture Cross-Check
- Experimental Validation Tasks
- Tasks
- Convention
 - Literature Reviews
 - Foundational (done)
 - Competing architectures (done — needed for related work)

- section)
 - Unread — lower priority, skip for preprint
- Architecture Derivation Writeups
- Experimental Validation Tasks
- Infrastructure
- Preprint Readiness Checklist
 - Ready now (no lab required)
 - Needs work before preprint (no lab)
 - Requires lab (post-preprint)
- Material Properties
 - ptr_glass
 - linbo3_tfln
 - sin_pic_al50
 - gaas_vcsl_850nm
 - ingaas_pin_detector
 - si_pin_detector_850nm
 - silica_fiber_smf28
 - gaas_vcsl_850nm_single_mode
- Design Parameters
 - optical
 - token_embedding
 - resonator
 - spatial
 - snr
 - interposer
 - model
 - power
- Economic Analysis (5-Year TCO)
- Performance Scenarios

Quantum Resonator Inference: System Design Document

Generated: 2026-04-28 **Status:** ARCH-1 through ARCH-17 LOCKED.
Experimental validation phase. **Repo:**
<https://github.com/jedelman/quantum-resonator-inference>

Quantum Resonator Inference — Architecture Specification

Status: ARCH-1 through ARCH-17 LOCKED **Last revised:** 2026-04-26 **Project:** All-optical resonator for embedded LLM token inference
Constraint: Single-tenant embedded device. No context switching.
No multi-tenancy. Offline training, static weights during inference.

1. Problem Statement

Standard LLM inference is memory-bandwidth-bound and thermodynamically inefficient. Every token requires moving billions of weight parameters through digital electronics at $\sim$ pJ/MAC. We are deriving an alternative: a coherent all-optical resonator that encodes model weights in the refractive index distribution of a physical medium and executes inference via photon-weight interactions at the speed of light.

The key question is not “how do we map a digital architecture to optics?” but “what computation does physics perform naturally, and can that computation be trained to implement token inference?”

The answer, derived below, is: **yes, and the natural primitive is a holographic resonator.**

2. Design Principles

1. **First-principles only.** Every architectural decision must be justified by physics, not by analogy to digital systems.
 2. **Coherence as a resource.** The design exploits optical coherence (phase, amplitude, spatial mode structure) as a computational degree of freedom.
 3. **Learning is structural.** Model weights are encoded in the refractive index distribution $\Delta n(x,y,z)$ of the resonator medium. Writing weights = writing holograms.
 4. **Token inference is wave dynamics.** Each forward pass is the time evolution of a wave field through a trained inhomogeneous medium. No digital approximation.
 5. **Embedded, single-tenant.** One model loaded once. No scheduling, no reloading, no context switching.
 6. **No undocumented parameters.** All values in parameters.toml require rationale.
 7. **In-situ training mandatory.** Weights cannot be transferred from simulation to physical cavity. Training must use the physical device as the forward model.
-

3. Theoretical Foundation (ARCH-1)

3.1 Wave Equation as RNN (Hughes et al. 2019)

The scalar wave equation for an optical field $u(x,y,t)$ in a medium with refractive index distribution $n(x,y)$:

$$\partial^2 u / \partial t^2 = (c_0 / n(x,y))^2 \cdot \nabla^2 u + f(x,y,t)$$

Finite-difference discretization with time step Δt yields:

$$u_{\{t+1\}} = 2u_t - u_{\{t-1\}} + \Delta t^2 (c_0 / n)^2 \nabla^2 u_t + \Delta t^2 f_t \quad (\text{eq. 1})$$

Define hidden state $h_t = [u_t, u_{\{t-1\}}]^T$. Then eq. 1 is structurally identical to an RNN update (Hughes et al. 2019, eq. 5):

$$h_t = A(n) \cdot h_{\{t-1\}} + P^{(i)} \cdot x_t \quad (\text{eq. 2})$$

$$y_t = |P^{(o)} \cdot h_t|^2 \quad (\text{eq. 3})$$

where $A(n)$ is determined by the refractive index distribution, $P^{(i)}$ is the input coupling matrix, $P^{(o)}$ is the readout matrix, and y_t is detected intensity. The trainable parameter is $n(x,y)$. Training = designing the refractive index distribution.

ARCH-1 conclusion: Any optical cavity with a spatially structured medium IS an RNN. The refractive index distribution is the weight matrix. Round-trip time = one RNN time step. This is an exact mapping at the level of Maxwell’s equations discretized in time, not an approximation.

3.2 Holographic Weight Storage (Psaltis et al. 1990)

A holographic grating in a photosensitive medium stores a weight matrix as a spatial modulation of the refractive index:

$$\Delta n(x, y) = \Delta n_{\max} \cdot \sum_k A_k \cdot \cos(k_k \cdot r + \phi_k) \quad (\text{eq. 4})$$

Each term encodes one outer product pattern. Angular multiplexing stores multiple patterns. For PTR glass: $\Delta n_{\max} = 5 \times 10^{-3}$ (Glebov 2010), thermally fixed after development, transparent at 850nm.

3.3 From Scalar Wave to Token Embedding

Token inference requires input $x \in \mathbb{R}^d$ and output $y \in \mathbb{R}^d$. Extension: excite the resonator with d simultaneous spatial modes, each amplitude-modulated by one component of the input token vector:

$$f(x, y, t) = \sum_{i=1}^d x_i \cdot \psi_i(x, y) \cdot \delta(t) \quad (\text{eq. 5})$$


After T round trips, the output field is read by d output detectors:

$$y_j = \left| \int \psi_j^*(x, y) \cdot u_T(x, y) \, dx \, dy \right|^2 \quad (\text{eq. 6})$$

The map $x \rightarrow y$ is a learned nonlinear transformation, optimized by training $\Delta n(x, y)$ via the adjoint method. Stacking multiple resonators (each with its own Δn) = stacking transformer layers.

4. System Architecture

4.1 System Overview



4.2 Resonator Geometry (ARCH-2, ARCH-3)

Geometry: Confocal Fabry-Perot, one flat mirror and one concave mirror ($R_c = L = 20\text{mm}$). PTR glass fills the cavity. HR mirror coatings deposited directly on PTR glass faces after holographic development — no separate optics inside the cavity.

The confocal condition ($R_c = L$) produces Hermite-Gaussian TEM_{mn} eigenmodes with fundamental waist $w_0 = \sqrt{\lambda L / 2\pi} = 5.2 \mu\text{m}$ at 850nm. Fresnel number $F = D^2 / (4\lambda L) = 78$ confirms >7000 modes fit within the 2.5mm aperture, of which 512 are addressed.

Why Fabry-Perot over ring: Ring resonators support traveling-wave modes, not standing-wave. The standing-wave mode structure of the Fabry-Perot naturally implements the $u_{\{t-1\}}$ term in eq. 1 via retroreflection. Ring geometry rejected.

Why PTR glass fills the cavity: The weight grating is a property of the glass bulk. Filling the cavity with glass rather than using a thin insert maximizes MVM interaction length (more Δn accumulated per pass). No air gaps inside the resonator.

ARCH-2/3 locked parameters:

Parameter	Value	Rationale
Cavity length L	20mm	Sets $\tau=133\text{ps}$, $T_{\text{coh}}=750 \gg$ $T_{\text{op}}=100$. Coherent regime.
Mirror reflectivity R	0.9990	Finesse=3140, intra-cavity gain $\sim 1000\times$, practical dielectric HR
Round-trip time τ	133.3 ps	$\tau = 2L/c_0$
Operating round trips T_{op}	100	$T_{\text{op}} \ll T_{\text{coh}}$. SNR $\geq 38\text{dB}$ at this depth.
Cavity aperture	2.5mm	Supports 512 modes at $50\mu\text{m}$ VCSEL pitch with $2\times$ margin
Modes addressed	512	Embedding dimension $d=512$
Polarization	Vertical linear only	VCSEL native, avoids mirror dichroism
VCSEL pitch	$50\mu\text{m}$	Glass Brain validated; $117\times$ safety vs $\lambda/2$
PTR glass geometry	$10\times 10\times 0.5\text{mm}$	Thermal management (ARCH-10); post-carve coating

4.3 The Pure-Glass Resonator

The inference resonator contains no active intra-cavity elements. The optical path during inference is:

VCSEL array (850nm, 512 modes)

- input face (HR coating $R=0.9990$, $\sim 0.1\%$ input coupling)
- PTR glass bulk ($\Delta n(x,y)$ weight grating)
- output face (HR coating, $\sim 10\%$ output coupling)
- Si PIN detector array

Nothing inside the resonator except glass. This is not a simplification for convenience — it is a physical requirement. Any active element (modulator, phase shifter, electrode) inside a Finesse-3140 cavity contributes insertion loss on every round trip. Even 0.1 dB/pass

accumulates to 10 dB over $T=100$ round trips, consuming the entire SNR budget. There is no margin for intra-cavity active elements at this finesse and round-trip count.

4.4 Input Encoding (ARCH-3)

Each token $x \in \mathbb{R}^{512}$ is encoded as a spatial amplitude distribution at the input mirror:

$$E_{in}(x,y) = \sum_{i=1}^{512} x_i \cdot \psi_i(x,y)$$

Implementation: 512 GaAs VCSELs at 850nm, arranged in a 23×23 grid at $50\mu\text{m}$ pitch (1.15mm footprint), collimated and magnified $2.2\times$ to fill the 2.5mm cavity aperture. Each VCSEL is amplitude-modulated at the token rate (75M tok/s = 13.3ns/token, requiring $\sim 75\text{MHz}$ modulation — well within VCSEL bandwidth of 5-10GHz).

Note: current encoding uses amplitude only (x_i real-valued $\rightarrow$ amplitude of mode ψ_i). Phase of each mode is unused at input. This is an unexploited degree of freedom. Phase encoding could double the effective input capacity. Deferred to future work; not required for initial validation.

4.5 Inter-Layer Coupling (ARCH-8)

Between successive resonator layers, coupling is entirely passive and free-space. No electronics, no active modulation, no homodyne reference in the inter-layer path.

Output face of layer k (partial transmission $\sim 10\%$)

- ↓ collimating lens
- ↓ relay lens pair (mode-match aperture $k \rightarrow$ aperture $k+1$)
- ↓ input face of layer $k+1$

The relay lens pair performs geometric mode-matching only — it maps the output field of cavity k onto the input aperture of cavity $k+1$ with correct magnification. It applies no computation. Inter-layer latency is $\sim 0.2\text{ns}$ for 50mm free-space propagation, negligible relative to the 13.3ns per-layer compute time.

Phase coherence is not required to be maintained across layers. The token embedding at each layer input is a real-valued intensity vector (the output of the Si PIN detector array), re-encoded as VCSEL amplitudes. Incoherent coupling between layers eliminates the need for a global phase reference, dramatically simplifying the multi-layer stack and enabling independent per-layer thermal management.

Inter-layer signal chain and activation function: The detector-driver-VCSEL chain implements the system's activation function. The signal path per mode is:

$P_k = E_k ^2$ per mode)	(layer k output intensity,
$I_{photo} = R \cdot P_k$ mA at $P=2.5\text{mW}$	$R = 0.6 \text{ A/W} \rightarrow I_{photo} \approx 1.5$
$V_{TIA} = R_f \cdot I_{photo}$ (within 3.3V rail)	$R_f = 667 \Omega \rightarrow V_{TIA} \leq 1\text{V}$
$I_{drive} = g \cdot V_{TIA}$ mA at nominal	$g = 5 \text{ mA/V} \rightarrow I_{drive} \approx 5$
$P_{VCSEL} = \eta_s \cdot \max(0, I_{drive} - I_{th})$ $\eta_s=0.6 \text{ W/A}$, $I_{th}=1 \text{ mA}$	

Combining: $P_{out} = \eta_s \cdot K \cdot \max(0, P_k - \theta)$ where: - $K = g \cdot R_f \cdot R = 2.0$ (dimensionless gain) - $\theta = I_{th} / K = 0.5 \text{ mW}$ (power threshold, tunable via VCSEL bias current) - $A^2 = \eta_s \cdot K = 1.2$ (net intensity

gain > 1, compensates coupling loss)

In the intensity (power) computational basis, this is exactly ReLU:

$$P_{\text{out}} = A^2 \cdot \max(0, P_{\text{in}} - \theta)$$

Linear MVMs operate on intensity vectors; ReLU on intensity is interleaved between every layer. This satisfies universal approximation (Leshno et al. 1993: any non-polynomial bounded activation suffices; ReLU is the degenerate limit $\theta \rightarrow 0$).

Key TIA design constraint: $R_f = 667 \, \Omega$, not $100 \, \text{k}\Omega$. At $I_{\text{photo}} = 1.5 \, \text{mA}$, $R_f = 100 \, \text{k}\Omega$ gives $V_{\text{TIA}} = 150 \, \text{V}$ — a rail violation. The correct transimpedance for $\leq 1 \, \text{V}$ output at this photocurrent is $667 \, \Omega$. This also corrects the SNR section which previously cited $100 \, \text{k}\Omega$ (that figure was from a low-light single-photon context, not applicable here).

Threshold tunability: VCSEL bias current sets the effective threshold: - Bias = 0: $\theta_{\text{eff}} = 0.5 \, \text{mW}$ (20% of P_{op}) — moderate nonlinearity - Bias = $95\% \cdot I_{\text{th}}$: $\theta_{\text{eff}} = 0.025 \, \text{mW}$ (1% of P_{op}) — near-linear (soft activation) - Bias < I_{th} always: VCSEL off when no signal — hard sparsity below threshold

The threshold is a single per-layer scalar set by bias DAC — trainable hyperparameter.

On the question of MZIs in the coupling path: MZI meshes are not present and are not needed. MZIs in photonic neural networks implement programmable weight matrices. In QRI, the weight matrix is implemented by the holographic grating inside the resonator. The coupling path has no computational role — it is purely geometric transport. Passive relay optics suffice.

4.6 Activation Function (ARCH-9, locked 2026-04-27)

The activation function is **ReLU on intensity**, implemented by the VCSEL threshold nonlinearity at each inter-layer boundary. No intra-cavity nonlinear optical element is required.

Derivation: See ARCH-8 signal chain. The complete transfer function from layer k intensity to layer $k+1$ intensity is:

$$P_{\text{out}} = A^2 \cdot \max(0, P_{\text{in}} - \theta)$$

where:

$$\begin{aligned} A^2 &= \eta_s \cdot K = \eta_s \cdot g \cdot R_f \cdot R = 1.2 && \text{(net gain)} \\ \theta &= I_{\text{th}} / K && = 0.5 \, \text{mW} \text{ (power threshold)} \\ K &= g \cdot R_f \cdot R && = 2.0 \end{aligned}$$

This is ReLU with gain $A^2 = 1.2$ and threshold $\theta = 0.5 \, \text{mW}$. The gain > 1 is intentional: it compensates the $\sim 10\%$ inter-layer coupling loss so that signal level is maintained across 24 layers without electronic amplification stages.

Proof of nonlinearity: $f(a \cdot P) = A^2 \cdot \max(0, a \cdot P - \theta) \neq a \cdot f(P)$ for $\theta \neq 0$. Fails homogeneity $\rightarrow$ nonlinear. ✓

Proof of expressiveness: ReLU networks with sufficient width are universal approximators (Hornik 1991). The intensity-domain computation (holographic MVM interleaved with intensity ReLU) is precisely this architecture. ✓

Why not Kerr SPM: $\varphi_{NL} \sim 10^{-15}$ rad/pass at operating intensity — negligible. SPM is a phase effect; it produces no amplitude nonlinearity at these power levels. Retired.

Computational basis note: The correct basis for QRI computation is **intensity** (power), not field amplitude. The holographic MVM mixes intensity modes; the VCSEL threshold applies ReLU to each intensity component; the next layer sees an intensity input. Phase is erased at each inter-layer boundary — this is intentional and required for thermal independence between layers (ARCH-8).

ARCH-9 locked parameters:

Parameter	Value	Derivation
Activation function	ReLU on intensity: $P_{out} = A^2 \cdot \max(0, P_{in} - \theta)$	VCSEL threshold
Net gain A^2	1.2	$\eta_s \cdot g \cdot R_f \cdot R = 0.6 \times 5e-3 \times 667 \times 0.6$
Power threshold θ	0.5 mW (zero bias)	$I_{th} / K = 1mA / 2.0$
θ tunability	0–0.5 mW via VCSEL bias DAC	Bias sets effective I_{th}
TIA transimpedance R_f	667 Ω	$V_{max}=1V$ at $I_{photo}=1.5mA$
Driver transconductance g	5 mA/V	$I_{drive,max}=5mA$ at $V=1V$
Intra-cavity nonlinear element	None	—
Cavity detuning	On-resonance ($\delta = 0$)	Maximize finesse buildup
Computational basis	Intensity (power)	Phase erased at each boundary
VCSEL bias stabilization	APC loop (automatic power control)	Standard driver IC feature; compensates I_{th} thermal drift (~3.8mA at 7.6K self-heating)
VCSEL PID lock bandwidth	~1 kHz	Thermal drift correction only

4.7 SNR and Noise Budget (ARCH-5)

Signal path:

Input: 2-3 mW per VCSEL mode

- finesse buildup ($\sim 1000\times$) → 2-3W intra-cavity
- $T=100$ round-trip loss: $0.995^{100} \approx 0.606$ (2.4 dB)
- output power ~1.2-1.8W aggregate
- Si PIN (0.6 A/W) → 0.36-0.54 A photocurrent
- TIA ($R_f=667\Omega$) → 6-bit quantizer

Dominant noise: shot noise. At $I \approx 0.4A$, readout bandwidth 1GHz:

$$\sigma_{shot} = \sqrt{(2eI\Delta f)} \approx 10\mu A$$

$$SNR = I/\sigma_{shot} \approx 40dB$$

Target for 6-bit precision: 38dB ($6.02 \times 6 + 1.76 = 37.9dB$). Margin: 2dB.

ARCH-5 locked parameters:

Parameter	Value
Input power per VCSEL	2-3mW
Intra-cavity power	2-3W
Output SNR	40dB (target ≥ 38 dB)
SNR margin	2dB
Noise floor	Shot-noise-limited

Risk: 2dB margin is tight. EXP-4 (thermal lensing) could erode it. Phase 1 VCSEL upgrade (+3dB, 2.5mW $\rightarrow$ 10mW) is available if margin is insufficient.

4.8 Token Throughput (ARCH-4)

$\tau = 2L/c_0 = 133.3 \text{ ps}$ (round-trip time)
 $t_{\text{token}} = T_{\text{op}} \times \tau = 100 \times 133.3 \text{ ps} = 13.3 \text{ ns}$
throughput = $1/t_{\text{token}} = 75\text{M tokens/sec}$

24-layer stack latency: $24 \times 13.3\text{ns} = 320\text{ns}$ optical + inter-layer propagation $\approx 330\text{ns}$ total. No electronic interposer latency in the inference path.

4.9 Holographic Weight Capacity (ARCH-7)

Weight matrix stored via low-rank factorization $W = U \cdot V^T$:

$U \in \mathbb{R}^{(512 \times r)}$, $V \in \mathbb{R}^{(512 \times r)}$
Total weights per layer: $1024r$

At rank $r=50$: 51.2k weights per layer. Each outer product $u_i \otimes v_i^T \rightarrow$ one holographic grating component, written via angular multiplexing. PTR glass supports ~ 1000 independent grating components (Glebov 2010, λ/D angular selectivity criterion) — ample headroom at rank-50.

For a 24-layer stack: $24 \times 51.2\text{k} = 1.23\text{M}$ equivalent parameters. This is a low-rank approximation; accuracy trade-off (rank-50 typically retains 95-98% of transformer capability) is acceptable given the physics constraints. Rank can be increased to 100-200 with mode basis expansion (ARCH-16).

4.10 Thermal Management (ARCH-10)

PTR plate geometry: $10 \times 10 \times 0.5\text{mm}$. Surface area 260mm^2 provides passive thermal dissipation with $\Delta T \approx 15\text{K}$ above ambient at 1W absorbed. Active Peltier cooling holds plate at $20\text{-}25^\circ\text{C}$ even at 10W intra-cavity (COP ~ 2 , $\sim 5\text{W}$ electrical overhead). Peltier is optional for the baseline design; included for margin.

5. Training Architecture

5.1 The Weight Translation Problem (ARCH-11)

Digital training — running the adjoint wave equation simulation on a GPU — computes an optimal $\Delta n(x,y)$ for a mathematical cavity. The physical cavity deviates from that model at scales that matter:

- A length error of $1\mu\text{m} = 1.2$ wavelengths at 850nm
- Mirror flatness errors of $\lambda/10$ introduce 85nm wavefront error per

- pass
- Internal stress gradients in the PTR blank produce spatially varying Δn not present in simulation

These deviations compound multiplicatively over $T=100$ round trips. The correspondence between simulated weights and physical weights is completely destroyed. **Weight translation from simulation to physical cavity is not feasible.** Every physical cavity is unique at the scale that determines the computation.

Training must therefore be performed in-situ, using the actual physical cavity as the forward model. The physical imperfections are automatically incorporated because gradients are computed from measurements made through the real glass.

5.2 Wavelength Separation: Write at 532nm, Infer at 850nm (ARCH-11)

Three isolation schemes for separating weight-writing from inference were analyzed.

Temporal separation (exploit slow grating decay between write and read phases) fails because any 850nm read beam has nonzero photorefractive cross-section σ_r in any optically responsive material. Over millions of inference tokens, the grating degrades via optical fixing fatigue. The erasure cannot be eliminated, only slowed. This provides no benefit over static PTR holography.

Spatial mode separation (write with TEM_{01} , read with TEM_{00} , exploit amplitude orthogonality) fails because amplitude orthogonality does not imply intensity orthogonality. The grating coupling coefficient involves an intensity overlap integral:

$$\kappa_{\{00,01\}} = (k/2n) \cdot \iint \psi_{00}^*(x,y) \cdot \Delta n_{01}(x,y) \cdot \psi_{00}(x,y) dx dy$$

Since $\Delta n_{01}(x,y)$ follows the TEM_{01} intensity profile (even function, nonzero overlap with Gaussian TEM_{00}), this integral is nonzero. The read beam both uses and partially erases the written grating. The computation mechanism and the erasure mechanism are identical and cannot be separated by spatial mode choice.

Wavelength separation (write at 532nm, read at 850nm) is categorically different. PTR glass photosensitivity is determined by the silver-cerium complex absorption spectrum, which peaks in the UV and falls to effectively zero at 850nm. The grating integrity condition is:

$$\sigma_r \cdot I_r \ll \sigma_w \cdot I_w$$

$$\sigma_r(850nm) \approx 0 \rightarrow \text{satisfied for any } I_r, \text{ unconditionally}$$

This is a physics argument — photons at 850nm lack the energy to drive the photochemical reaction that creates or destroys the grating — not a rate argument or an overlap argument. The inference beam cannot erase the weights. This property does not degrade with operating time, temperature, or intensity within the PTR stability envelope.

Wavelength separation is locked as the isolation mechanism.

5.3 In-Situ Training Protocol (ARCH-11)

Training uses the physical inference cavity as the forward model. The mechanical housing includes a 532nm write port (capped during inference) in addition to the 850nm input port.

Forward pass (850nm): Inject training input x via the VCSEL array. Circulate $T=100$ round trips. Measure output y via the Si PIN detector array. Compute loss $L = \|y - y_{\text{target}}\|^2$.

Gradient computation (digital): The adjoint method computes $\partial L / \partial \Delta n(x,y)$ on a GPU using the measured output as the boundary condition. The physical cavity's imperfections are implicitly absorbed because the forward pass is measured, not simulated.

Weight update (532nm): The gradient $\partial L / \partial \Delta n(x,y)$ is encoded as a spatial amplitude and phase pattern on the 532nm write beam, which is injected via the write port in single-pass configuration. The 532nm beam's intensity pattern drives a photorefractive index increment matching the gradient step:

$$\Delta n(x,y) \leftarrow \Delta n(x,y) + \eta \cdot \partial L / \partial \Delta n(x,y)$$

where η is set by the 532nm exposure dose. The 850nm beam is off during writing.

Batch gradient accumulation: Individual training samples should not trigger individual 532nm write pulses. Consecutive write exposures create overlapping latent images that partially erase each other (holographic crosstalk), degrading the weight trajectory. Instead, gradients are accumulated digitally over a batch of training tokens and a single summed gradient is written per batch. This maps onto standard batch gradient descent and eliminates intra-batch crosstalk. Batch size is a free parameter; larger batches reduce crosstalk and improve gradient quality at the cost of slower feedback. Optimal batch size is determined by the ratio of individual gradient magnitude to accumulated grating strength, which EXP-7 will characterize.

Thermal development: After one full training epoch (all batches processed), the PTR glass is removed from the cavity and thermally developed (500°C, ~30 minutes). Development converts the latent photorefractive exposure (silver nanoparticle seeds, ~10-20% of final Δn) into permanent crystallographic index modulation (NaF nanoparticle growth, full Δn). The crystallographic grating is immune to subsequent 532nm exposures — it cannot be overwritten or erased. Subsequent write cycles add corrections on top of the fixed structure.

Iteration: Reinstall the developed glass, remeasure the forward pass, evaluate residual loss. If loss is above target, run another write-develop cycle. Convergence expected in 3-5 cycles. The system identification loop (see ARCH-13) progressively refines the adjoint simulation to match the physical cavity, reducing residual error each cycle.

Training timeline (per layer):

Step	Duration	Notes
One write epoch	~1 hour	532nm power and target Δn magnitude; EXP-3
Thermal development	30 min	Batch all 24 layers simultaneously
Reinstall + evaluate	15 min	Kinematic mount; no realignment

Cycles to convergence	3-5	EXP-7 validates
Total, 24 layers parallel	~1 day	One furnace batch, 24 write stations

5.4 Grating Stability During Inference (ARCH-11)

After thermal development, the crystallographic grating is permanent. The silver-cerium photosensitive complex has been consumed; subsequent 850nm photons find no reactive species. Glebov et al. demonstrated PTR grating stability over decade-plus timescales with no measurable diffraction efficiency change. The inference beam at 2-3W intra-cavity poses no grating integrity risk.

The only inference-phase stability concern is thermal drift of the effective optical path length, which shifts the cavity resonance frequency (not the grating). This is addressed by the VCSEL PID frequency lock (ARCH-9) and flagged for validation as EXP-4.

5.5 Batch Size and Epoch Scale (ARCH-11)

Because batch gradients are accumulated digitally before each 532nm write, the effective batch size can be arbitrarily large — limited only by GPU memory for gradient accumulation, not by any optical constraint. At 75M tok/s, running a batch of 1 billion training tokens takes ~13 seconds of optical forward-pass time. This enables training on massive corpora per write-develop cycle, which is a significant advantage over any digital training setup: the optical forward pass is essentially free relative to the gradient computation. Training epochs can be made as large as desired; the per-epoch cost is dominated by the 30-minute furnace cycle, not by the optical exposure time.

6. Scaling Architecture

6.1 Block Duplication (ARCH-17)

A single QRI block (one PTR glass resonator + relay optics) implements one transformer layer. Blocks can be duplicated in two configurations.

Series (depth): Stack N blocks, each with its own holographic weight plate. The output of block k feeds the input of block k+1 via passive relay optics. Computation depth scales with N. This is the 24-layer baseline.

Parallel (width and expertise): Route the same input to M independent blocks simultaneously. Each block computes a different linear transformation (different holographic weight plate) over a different subspace of the embedding. Outputs are concatenated or combined. This is the mixture-of-experts (MoE) architecture, implemented physically as an array of independent glass resonators with passive optical splitters and combiners. M can be large because each block is independent — no shared optical resource creates a bottleneck.

Incoherent coupling between layers and between parallel blocks is required for scalability. If phase coherence were maintained across blocks, a global phase reference would be needed for every block, and thermal drift in any one block would corrupt all others. Incoherent

coupling (intensity readout at each block output, re-encoding at each block input) makes each block thermally and mechanically independent. This is the correct design for any system with more than a few blocks.

6.2 Clone-and-Fine-Tune for Scale Training (ARCH-17)

The weight translation problem (§5.1) establishes that weights cannot be transferred between a simulation and a physical cavity. The question for scaling is whether weights can be transferred between two physical cavities — i.e., whether unit 02 can be initialized from unit 01's trained grating and then fine-tuned, rather than trained from scratch.

Why naive cloning fails in general. The trained grating of unit 01 encodes two superimposed components that cannot be separated by inspection of the physical glass: the intended computation (the weight matrix W that the training procedure converged to) and the cavity-specific correction (the Δn pattern that compensates for unit 01's specific mirror figure errors, glass inhomogeneities, and length deviations). When this grating is copied to unit 02, the cavity-specific correction terms — which were tuned to unit 01's imperfections — are wrong for unit 02's different imperfections. They introduce errors rather than correcting them.

Why cloning still works as an initialization. Despite this, the cloned grating is not useless for unit 02. The intended computation W is present, even if distorted. Unit 02 is computing an imperfect version of the right computation, not a random one. This is precisely the situation where fine-tuning is effective: the loss landscape has already been explored and a good region found; fine-tuning is a local search to correct for the cavity mismatch, not a global search from scratch.

The tolerance condition. For clone-and-fine-tune to be cheaper than full training, the following must hold: the cavity-specific correction $|\Delta n_{\text{comp2}} - \Delta n_{\text{comp1}}|$ (the difference between unit 02's required correction and unit 01's baked-in correction) must be small relative to the total grating magnitude $|\Delta n_{\text{total}}|$. This difference is bounded by the inter-unit manufacturing variation — how consistently the two cavities are built. Manufacturing consistency is therefore a direct determinant of the clone-and-fine-tune efficiency.

The two-stage training economy. If manufacturing consistency is sufficient that fine-tuning converges in 1-2 write-develop cycles rather than the 3-5 needed for full training, the scaling economics change dramatically. A factory of N inference units requires: one fully-trained master unit (full training, ~ 1 day), an optical grating copying apparatus (contact printing or holographic duplication), and per-unit fine-tuning (1-2 write-develop cycles, ~ 4 -8 hours each). The fine-tuning cost is sublinear in total training cost per unit. This maps exactly onto the digital ML paradigm of foundation model pretraining followed by deployment-specific fine-tuning — implemented in glass.

What must be validated. The clone-and-fine-tune strategy requires experimental validation before being relied upon for scale deployment. EXP-7 (see §7) is extended to include this validation.

6.3 Deployment Model (ARCH-12, ARCH-13)

During inference, the device operates as a pure 850nm system. The 532nm write port is physically capped. No training, gradient computation, or weight updates occur during inference.

The single-cavity phase stability budget is approximately:

$$\sigma_{\text{per_token}} \sim 1\text{-}5 \text{ mrad/token (estimated, EXP-5 validates)}$$
$$B = (\pi/4)^2 / \sigma^2_{\text{per_token}} \approx 24,600 \text{ tokens at } \sigma=5 \text{ mrad/token}$$

This is adequate for LLM inference context lengths of 2K-128K tokens without ensemble averaging. Multi-cavity ensemble arrays are a scaling option if longer context or higher phase precision is required, but are not part of the baseline design.

Model update = glass swap. The device chassis is permanent hardware. PTR glass plates (10×10×0.5mm, ~\$10-50 per blank) are swappable via kinematic mounts in minutes. No power cycling, no software update, no re-initialization. The device serves tokens immediately when the 850nm laser is on and the cavity is locked.

Write station. Training is performed at a dedicated write facility, not on the deployed device. The write station consists of a 532nm CW laser (Nd:YAG SHG, ~100mW), an SLM or holographic beam shaper for gradient encoding, a kinematic stage, and a reference beam path. The write station couples to the inference cavity via the write port, or the glass plate can be written on a standalone holographic bench. Finished plates are shipped to the deployed device. A model update is operationally equivalent to receiving a small package of glass cards.

Adjoint simulation infrastructure. Gradient computation ($\partial L / \partial \Delta n(x,y)$ for each layer) runs on a GPU cluster using the differentiable wave equation model. The simulation is coupled to physical measurement in a closed loop: simulate → write → measure → compare → update simulation parameters → repeat. Over multiple cycles, the simulation converges to an accurate surrogate for the physical cavity (system identification). By the final write cycle, the surrogate may be good enough to compute updates without further in-situ measurement.

7. Rank Scaling (ARCH-16)

7.1 Rank and Mode Basis

RNN state dimensionality d = rank of the learned weight tensor. Each spatial mode carries one state dimension. The holographic grating encodes the weight matrix as a sum of rank-1 outer products, one per grating component. The mode basis determines how many independent grating components can be supported before the cavity Q degrades.

Hermite-Gauss (default, rank <100): Orthogonal in 2D rectangular geometry. 10×10 grid = 100 modes at rank-50; 20×20 grid = 400 modes at rank-100 (limited by diffraction loss above ~rank-100).

Laguerre-Gauss (rank >100): Cylindrical symmetry basis. Radial index p + azimuthal index ℓ . Scales more efficiently to high rank; mode overlap loss lower above rank-100.

Hybrid basis (recommendation for rank >100):

$$\Phi_{\text{hybrid}} = \{\text{HG}_{\{m,n\}} : m+n < 10\} \cup \{\text{LG}_{\{p,\ell\}} : p \leq 15, \ell \leq 3\}$$
$$\approx 100 \text{ HG} + 100 \text{ LG} = \text{rank-200 native capacity}$$

7.2 Rank Scaling Regimes

Regime 1: Rank 50-100 (safe, current design) Insertion loss ~2.2-3.0dB, SNR 38-40dB, thermal PID at 1kHz sufficient. Accuracy within 1-2% of digital baseline.

Regime 2: Rank 100-200 (pushing limits) Insertion loss ~3.5-4.5dB, SNR 36-37dB, mode coupling introduces spurious terms. Requires mode orthonormalization regularization (λ_3 penalty on $\langle \phi_i | \phi_j \rangle^2$). Expected 2-4% accuracy loss vs. digital baseline.

Regime 3: Rank >200 (failure) Cavity Q degrades below critical threshold. SNR drops to ~30dB. Thermal coupling between adjacent cavity regions exceeds PID tracking bandwidth. Not recommended.

7.3 Rank-Loss Tradeoff

Test accuracy = Digital_baseline - ϵ_{rank} - ϵ_{SNR}

$\epsilon_{\text{rank}} \approx 1\% \cdot \log_2(\text{rank}) / \log_2(256)$
 $\epsilon_{\text{SNR}} \approx 0.5\% \cdot (40\text{dB} - \text{SNR}_{\text{actual}}) / 10\text{dB}$

Rank	SNR (dB)	ϵ_{rank} (%)	ϵ_{SNR} (%)	Total loss (%)
50	37.8	0.8	0.1	0.9
100	37.0	1.5	0.5	2.0
150	36.2	2.1	0.9	3.0
200	35.5	2.8	1.3	4.1

Production target: Rank-100 (200 basis modes, ~2% total accuracy loss, within acceptable range). **Stretch goal:** Rank-150 (requires +3dB VCSEL upgrade for margin). **Ceiling:** Rank-200 (all margins exhausted; not recommended for production).

Note: the ϵ_{rank} and ϵ_{SNR} model accounts for MVM fidelity and SNR. The activation function (ReLU on intensity, VCSEL threshold) is guaranteed by electronics and does not contribute accuracy loss.

8. Convergence Theory (ARCH-14, ARCH-15)

8.1 Photonic Backpropagation (ARCH-14)

Hughes et al. (2018) proved photonic backpropagation exactness via adjoint variable methods. For a photonic linear transformation $y = M \cdot x$, gradients with respect to structural parameters are:

$\partial L / \partial \phi = \text{Re} (\partial L / \partial y \mid \partial y / \partial \phi)$

computed by time-reversing the optical field and interfering with the loss gradient signal. This is exact — not an approximation — for any linear photonic system.

In QRI’s in-situ training protocol, the adjoint computation runs digitally using the measured forward pass output as the boundary condition. The physical cavity’s imperfections are automatically captured. Pai et al. (2023) demonstrated this experimentally, achieving ~94% MNIST accuracy (matching digital training) using in-situ photonic backpropagation.

Convergence rate bound (Hughes 2018, implied):

$$E[||\nabla L||^2] \leq O(1/\sqrt{N_{\text{circ}}}) \quad \text{where } N_{\text{circ}} = \text{round trips per forward pass}$$

At $N_{\text{circ}} = 100$: gradient noise $\sim 1\%$ of loss signal. Learning rate 0.01-0.1 per epoch. Expected convergence: 10-100 epochs (hours of optical forward-pass time).

8.2 Loss Landscape (ARCH-15)

The photonic loss (intensity measurement) and the digital loss (cross-entropy or MSE) are related but not identical:

$$L_{\text{optical}} = \int |E_{\text{out}}(t) - E_{\text{target}}(t)|^2 dt$$
$$L_{\text{digital}} = \sum_i ||y_{\text{pred}_i} - y_{\text{target}_i}||^2$$

$L_{\text{optical}} \approx L_{\text{digital}}$ when amplitude matches normalized prediction and target phase encoding tracks cavity mode phase. Phase mismatch can create spurious minima in L_{optical} not present in L_{digital} . Mitigation: pre-calibrate target phase encoding to track cavity resonance condition throughout training.

Training stability requires signal-to-gradient ratio $>10:1$ (satisfied at 40dB SNR: gradient noise $\sim 1\%$), learning rate <0.1 per epoch, and batch size $\gg 1$ (batch accumulation before each write, as specified in §5.3).

9. Decision Log

Date	Decision	Rationale
2026-04-19	ARCH-1 LOCKED: Fabry-Perot cavity as RNN	Hughes 2019 exact mapping: wave eq = RNN update
2026-04-19	ARCH-1 LOCKED: PTR glass as weight medium	Psaltis 1990 + Glass Brain validation. Non- volatile.
2026-04-19	ARCH-1 LOCKED: 850nm wavelength	GaAs VCSEL OTS maturity, Si PD response, PTR transparency. Glass Brain continuity.
2026-04-19	Embedded/single- tenant constraint	No HBM3, no scheduling. One model, static weights.
2026-04-20	ARCH-2 LOCKED	L=20mm, R=0.9990, $\tau=133.3\text{ps}$, $T_{\text{op}}=100$. Coherent regime.
2026-04-20	ARCH-3 LOCKED	512 spatial modes, 2.5mm aperture, single vertical polarization.
2026-04-20	ARCH-4 LOCKED	75M tok/s (13.3ns/token).
2026-04-20	ARCH-5 LOCKED	SNR 40dB $\geq$ 38dB target. Shot-noise- limited.
2026-04-20	ARCH-7 LOCKED	Rank-50 $U \cdot V^T$, 51.2k weights/layer, 1.23M

		total.
2026-04-20	ARCH-8 LOCKED	All-optical layer coupling, passive relay optics only. No MZIs in coupling path.
2026-04-20	ARCH-9 LOCKED	Kerr SPM (χ^3 in PTR glass), $\phi_{NL}=0.2-1$ rad/pass.
2026-04-27	ARCH-9 REVISED	Kerr SPM retired.
2026-04-27	ARCH-9 LOCKED	Activation: ReLU on intensity $P_{out}=A^2 \cdot \max(0, P_{in}-\theta)$. $A^2=1.2$, $\theta=0.5\text{mW}$. TIA $R_f=667\Omega$ (corrected from $100k\Omega$). Computational basis: intensity.
2026-04-20	ARCH-10 LOCKED	$10 \times 10 \times 0.5\text{mm}$ plate, passive+active thermal management.
2026-04-24	ARCH-14 LOCKED	Photonic backprop exactness via adjoint (Hughes 2018, Pai 2023).
2026-04-24	ARCH-15 LOCKED	Loss landscape: $L_{optical} \approx L_{digital}$ with phase calibration.
2026-04-26	ARCH-16 LOCKED	Rank ceiling 200, production target rank-100. Hybrid HG/LG basis above rank-100.
2026-04-26	ARCH-11 RETRACTED	Ephemeral weights via intra-cavity LiNbO3 MZM: disqualifying. $0.1 \text{ dB/pass} \times T=100 = 10 \text{ dB}$ cumulative insertion loss per token. Wipes SNR budget.
2026-04-26	ARCH-11 (new) LOCKED	Weight translation infeasible (sub- λ manufacturing error $\times T=100$). In-situ training mandatory. Wavelength separation (532nm write, 850nm infer) is the only clean isolation mechanism — $\sigma_r(850\text{nm}) \approx 0$ is a physics argument, not a rate tradeoff. Pure-glass resonator. Batch gradient accumulation before each write. Thermal cure between epochs.
2026-04-26	ARCH-12 REVISED	Ensemble cavity arrays demoted to scaling option. Single-cavity phase budget ($\sim 24K$

2026-04-26	ARCH-13 REVISED	tokens) adequate for LLM inference. Model update = glass swap. Write station at dedicated facility. Adjoint simulation + system identification loop.
2026-04-26	ARCH-17 LOCKED	Clone-and-fine-tune scaling strategy. Grating encodes W + cavity-specific correction (inseparable). Cloning transfers W approximately; fine-tuning corrects cavity mismatch. Efficiency conditional on manufacturing consistency. EXP-7 extended to validate.
2026-04-26	EXP-6 CLOSED	LiNbO3 MZM removed from design. Single-pass insertion loss measurement no longer relevant.

10. Open Experiments

ID	Priority	Description	Blocks
EXP-1	CLOSED 2026-04-27	PTR χ^3 (n_2) at 850nm — no longer required. Kerr SPM retired. Activation function is detector-squaring via VCSEL driver.	—
EXP-2	HIGH	Two-wavelength photosensitivity: PTR under 532nm write + 850nm read simultaneously. Confirm $\sigma_r(850nm) \approx 0$ empirically.	ARCH-11 isolation claim
EXP-3	HIGH	Grating growth rate: Δn vs. 532nm exposure time at accessible intensities. Sets write epoch duration. Thermal	ARCH-11 timeline

EXP-4	HIGH	lensing: cavity stability under 2-3W CW 850nm intra-cavity. Acceptable drift <5 mrad/hour.	ARCH-5 SNR margin
EXP-5	MED	Homodyne phase-lock stability: VCSEL PID lock over 1-hour inference run.	ARCH-12 phase budget
EXP-6	CLOSED	LiNbO3 MZM insertion loss at 850nm. Closed 2026-04-26: MZM removed from design.	—
EXP-7	HIGH	In-situ training convergence + clone-and-fine-tune validation. See §10.1 below.	ARCH-11, ARCH-17

10.1 EXP-7: In-Situ Training Convergence and Clone Validation (Expanded)

Objective: Validate the two core claims of the training architecture — that in-situ iterative write-develop training converges in a tractable number of cycles, and that clone-and-fine-tune is viable as a scaling strategy.

Phase A — Convergence rate. Build a minimal single-layer holographic RNN at rank-10. Train from a blank PTR glass plate using the in-situ two-wavelength protocol with batch gradient accumulation and iterative thermal development. Measure loss vs. write-develop cycle number. Target: convergence to within 2% of the digital baseline (simulated rank-10 RNN on the same task) in ≤ 5 cycles. If convergence is slower, identify the dominant error source from the following candidates: (a) gradient encoding fidelity — does the 532nm write pattern accurately reproduce the computed $\partial L / \partial \Delta n$? (b) thermal development precision — does the furnace cycle faithfully convert latent image to crystallographic grating without distortion? (c) cavity reinstallation error — does the kinematic mount reproduce the cavity mode structure to sufficient precision?

Phase B — Clone-and-fine-tune. After Phase A produces a converged unit 01, optically copy the trained grating to a fresh PTR blank using contact holographic printing (or equivalent duplication method). Install the clone in a second, independently-built cavity (unit 02) with nominally identical but physically distinct specifications. Measure the initial loss of unit 02 with the cloned grating (no fine-tuning). Then run the fine-tuning protocol (iterative write-develop cycles starting from the cloned grating) and measure loss vs. cycle number. Compare the number of fine-tuning cycles required for unit 02 to the number of full training cycles required for unit 01.

Success criteria: - Phase A: ≤ 5 write-develop cycles to within 2% of digital baseline. - Phase B: Clone initializes unit 02 at lower loss than blank glass. Fine-tuning converges in ≤ 2 cycles (vs. ≤ 5 for full training). Confirms that clone-and-fine-tune is meaningfully cheaper than full training.

Secondary measurements: - Batch size sensitivity: run Phase A at batch sizes of 100, 1000, 10000 tokens per write. Characterize crosstalk degradation vs. batch size. Identify the minimum batch size at which training is stable. - Manufacturing consistency: measure cavity length, mirror figure, and glass homogeneity of unit 01 and unit 02. Correlate inter-unit physical differences with the initial loss of the cloned unit 02 and the number of fine-tuning cycles required. This establishes the manufacturing tolerance requirement for clone-and-fine-tune to be efficient.

11. Architecture Summary

ARCH	Component	Status	Key Result
1	Wave eq $\rightarrow$ RNN	✓ LOCKED	Hughes 2019 exact mapping. Resonator IS an RNN.
2	Cavity geometry	✓ LOCKED	L=20mm, R=0.9990, Finesse=3140, τ =133ps, T _{op} =100
3	Mode structure	✓ LOCKED	512 HG modes, 2.5mm aperture, 50 μ m VCSEL pitch
4	Throughput	✓ LOCKED	75M tok/s, 13.3ns/token
5	SNR budget	✓ LOCKED	40dB achieved, 38dB required, 2dB margin
6	Training (old)	✗ SUPERSEDED	Replaced by ARCH-11
7	Weight capacity	✓ LOCKED	Rank-50, 51.2k weights/layer, 1.23M total
8	Layer coupling	✓ LOCKED	Passive relay optics. No MZIs. Incoherent between layers.
9	Activation function	✓ LOCKED 2026-04-27	ReLU on intensity: $P_{out}=A^2 \cdot \max(0, P_{in}-\theta)$. $A^2=1.2$, $\theta=0.5\text{mW}$. VCSEL threshold. Intensity computational basis.
10	Thermal management	✓ LOCKED	10 $\times$ 10 $\times$ 0.5mm plate, Peltier optional
11	In-situ training	✓ LOCKED (2026-04-26)	Pure-glass resonator. 532nm write / 850nm infer. Batch accumulation + thermal cure per epoch. Weight translation infeasible; training must be in-situ.
12	Deployment	✓ LOCKED (2026-04-26)	Single cavity. Glass swap = model update. $\sim 24\text{K}$ token phase budget.

13	Training infra	✓ LOCKED (2026-04-26)	Write station + adjoint simulation. System identification loop. ~1 day/update.
14	Convergence proof	✓ LOCKED	Adjoint backprop exact (Hughes 2018). Empirical: 94% MNIST (Pai 2023).
15	Loss landscape	✓ LOCKED	$L_{\text{optical}} \approx L_{\text{digital}}$ with phase calibration. Batch size >1 required.
16	Rank scaling	✓ LOCKED	Production rank-100, ceiling rank-200. Hybrid HG/LG basis.
17	Clone scaling	✓ LOCKED (2026-04-26)	Clone-and-fine-tune viable conditional on manufacturing consistency. EXP-7B validates.

Architecture Cross-Check

All architecture decisions validated for mutual consistency by
analyze/arch_crosscheck.py.

=====

ARCHITECTURE CROSS-CHECK: ARCH-1 through ARCH-10

=====

ARCH-1 ↔ ARCH-2	✓ PASS	ARCH-1 ↔ ARCH-2: Resonator geometry supports wave RNN primitive
ARCH-2 ↔ ARCH-3	✓ PASS	ARCH-2 ↔ ARCH-3: Geometry supports 512 modes with 14.5× margin
ARCH-3 ↔ ARCH-4	✓ PASS	ARCH-3 ↔ ARCH-4: 512 modes transport 75M tok/s with 13.3ns latency
ARCH-4 ↔ ARCH-5	✓ PASS	ARCH-4 ↔ ARCH-5: 75M tok/s achieves 40dB SNR with 2.5mW input
ARCH-5 ↔ ARCH-6	✓ PASS	ARCH-5 ↔ ARCH-6: 40dB SNR supports coherent Hebbian via 532nm
ARCH-6 ↔ ARCH-7	✓ PASS	ARCH-6 ↔ ARCH-7: Rank-50 factorization (1.23M params) fits in 1000 multiplexed gratings
ARCH-7 ↔ ARCH-8	✓ PASS	ARCH-7 ↔ ARCH-8: Holographic weights readable via Homodyne detection
ARCH-8 ↔ ARCH-9	✓ PASS	ARCH-8 ↔ ARCH-9: ReLU on intensity. $A^2=1.20$, $\theta=0.5\text{mW}$, $V_{\text{TIA}}=1000\text{mV}$ ✓
ARCH-9 ↔ ARCH-10	✓ PASS	ARCH-9 ↔ ARCH-10: VCSEL external to cavity (15K rise is glass, not VCSEL). VCSEL self-heating ~7.6K → $\Delta I_{\text{th}} \sim 3.8\text{mA}$, compensated by APC loop.
Full Integration	✓ PASS	ARCH-1 through ARCH-10 fully integrated and mutually consistent
NARG Compat	✓ PASS	NARG: disabled (baseline mode)
PTYCH Compat	✓ PASS	PTYCHOGRAPHY: disabled (standard holography)
NARG+PTYCH	✓ PASS	NARG+PTYCH: not both enabled (no interaction)

=====

RESULT: All architecture decisions LOCKED and mutually consistent.
Ready for experimental validation phase (EXP-2 through EXP-5, EXP-7).



Experimental Validation Tasks

Tasks

Convention

- [] open
 - [x] done
 - Priority: HIGH / MED / LOW
-

Literature Reviews

Foundational (done)

- ☒ Hughes et al. 2019, Sci. Adv. — Wave physics as analog RNN. Core mapping: wave eq = RNN update. Load-bearing for ARCH-1. See [citations/hughes_2019_wave_rnn.md](#).
- ☒ Psaltis et al. 1990, Nature — Holography in artificial neural networks. Hologram = weight matrix. Load-bearing for ARCH-6/7. See [citations/psaltis_1990_holography_nns.md](#).
- ☒ Fu et al. 2024, Light: Sci. Appl. — ONN review. Field map. See [citations/onn_review_2024.md](#).
- ☒ Vicentini et al. 2021, Nat. Photonics — Dual-comb holography. See [citations/dual_comb_holography_2021.md](#).

Competing architectures (done — needed for related work section)

- ☒ Lin et al. 2018, Science — D²NN. Feedforward diffractive ONN, fixed weights. See [citations/lin_2018_d2nn.md](#).
- ☒ Shen et al. 2017, Nat. Photonics — MZI mesh ONN, unitary constraint. See [citations/shen_2017_mzi_mesh.md](#).
- ☒ Feldmann et al. 2019, Nature — PCM spiking ONN. Endurance-limited. See [citations/feldmann_2019_pcm_snn.md](#).
- ☒ Feldmann et al. 2021, Nature — Photonic tensor core (WDM + PCM). See [citations/feldmann_2021_tensor_core.md](#).
- ☒ Zhong et al. 2023, Nat. Commun. — Graphene/Si activation function. Considered and rejected (intra-cavity loss). See [citations/zhong_2023_graphene_activation.md](#).
- ☒ Duport et al. 2012, Opt. Express — All-optical reservoir computing. SOA delay loop. See [citations/duport_2012_reservoir.md](#).
- ☒ Xu et al. 2021, Nature — 11 TOPS photonic convolutional accelerator. See [citations/xu_2021_tops_comb.md](#).
- ☒ Xu et al. 2024, Science — Taichi 160 TOPS/W photonic chiplet. Benchmark to address. See [citations/xu_2024_taichi.md](#).
- ☒ Farhat et al. 1985, Appl. Opt. — First optical neural network (Hopfield). Historical anchor. See [citations/farhat_1985_hopfield.md](#).
- ☒ Reck et al. 1994, Phys. Rev. Lett. — Unitary decomposition via MZI

mesh. QRI explicitly avoids unitary constraint. See
citations/reck_1994_unitary.md.

Unread — lower priority, skip for preprint

- ☐ Bogaerts et al. 2020, Nature — Programmable photonic circuits. (MED — general PIC state of art, not directly competing)
 - ☐ Picqué & Hänsch 2019, Nat. Photonics — Frequency comb spectroscopy review. (LOW — dual-comb encoding deferred)
 - ☐ Shams-Ansari et al. 2020 — TFLN dual-comb. (LOW — dual-comb encoding deferred)
 - ☐ Coddington et al. 2009, Nat. Photonics — Dual-comb ranging. (LOW)
 - ☐ Ideguchi et al. 2013, Nature — Coherent Raman dual-comb. (LOW)
-

Architecture Derivation Writeups

These decisions are LOCKED in architecture.md. These tasks are formal derivation documents for the paper's methods section.

- ☐ ARCH-2 derivation: Resonator geometry (Fabry-Perot). Formally derive $L=20\text{mm}$, $R=0.9990$, $T_{\text{op}}=100$ from wave RNN coherence requirement and SNR budget. Currently locked but derivation is inline in architecture.md — needs standalone derivation doc.
 - ☐ ARCH-3 derivation: Mode structure. Formally derive 512 modes fit in 2.5mm aperture at $50\mu\text{m}$ pitch; Fresnel number; Hermite-Gaussian mode orthogonality.
 - ☐ ARCH-4 derivation: Token throughput. Formal derivation of 75M tok/s from $\tau=133\text{ps}$ and $T_{\text{op}}=100$.
 - ☐ ARCH-5 derivation: SNR budget. Formal shot-noise derivation to 40dB; 6-bit precision requirement.
 - ☐ ARCH-6 derivation: Training pipeline. Adjoint method for wave dynamics; $\Delta n(x,y)$ update rule; 532nm write protocol.
 - ☐ ARCH-7 derivation: Hologram capacity. Angular multiplexing; rank-50 fits in PTR plate at $50\mu\text{m}$ grating pitch; parameter count.
 - ☐ ARCH-8 derivation: Inter-layer coupling. Relay lens pair geometry; incoherent coupling justification; VCSEL driver signal chain.
-

Experimental Validation Tasks

All require lab. Do not block preprint on these.

- ☒ EXP-1 (CLOSED 2026-04-27): PTR χ^3 @ 850nm — CLOSED. Kerr SPM retired. Activation is VCSEL threshold ReLU.
- ☐ EXP-2 (HIGH): Two-wavelength photosensitivity — PTR @ 532nm write + 850nm read simultaneously. Confirm $\sigma_r(850\text{nm})\approx 0$ (no cross-sensitization). Blocks ARCH-11 isolation claim.
- ☐ EXP-3 (HIGH): Hebbian grating growth rate — Δn vs. 532nm exposure time. Target: $\Delta n=5\times 10^{-3}$ in <1000 inference passes. Sets write epoch duration.
- ☐ EXP-4 (HIGH): Thermal lensing dn/dT — cavity stability under 2-3W CW intra-cavity load. Acceptable drift: <5 mrad/hour. Blocks SNR margin confidence.
- ☐ EXP-5 (MED): Homodyne phase-lock stability — VCSEL PID lock over 1-hour inference run. Blocks ARCH-12 phase budget.
- ☒ EXP-6 (CLOSED 2026-04-26): LiNbO_3 MZM insertion loss — CLOSED. MZM removed from design.
- ☐ EXP-7 (HIGH): In-situ training convergence. Phase A: rank-10, single layer, ≤ 5 write-develop cycles to 2% of digital baseline.

Phase B: clone-and-fine-tune (unit 01 → unit 02, ≤ 2 cycles). Blocks ARCH-17 validation.

Infrastructure

- ☐ INFRA-1: PDF fetching in generate_sysdoc.py — download PDFs where DOI available, store in citations/
 - ☐ INFRA-2: design/render_resonator.py — resonator geometry diagram (confocal Fabry-Perot, mode structure, PTR plate, VCSEL array)
 - ☐ INFRA-3: Conversations/ log rotation — one file per session (currently done manually)
-

Preprint Readiness Checklist

Ready now (no lab required)

- ☒ Architecture: ARCH-1 through ARCH-17 locked
- ☒ Crosscheck: all 13 checks PASS
- ☒ Literature: all major competing works reviewed, differentiation documented
- ☒ Signal chain: fully derived (TIA, driver, VCSEL, activation)
- ☒ Activation function: ReLU on intensity, proven nonlinear, universal approximation
- ☒ Economic analysis: 5-year TCO vs. GPU
- ☒ System doc: auto-generated PDF from current sources

Needs work before preprint (no lab)

- ☐ ARCH-2 through ARCH-8 formal derivation docs (methods section material)
- ☐ Related work section: differentiation table vs. D²NN, MZI, PCM, reservoir, Taichi
- ☐ Abstract + introduction draft
- ☐ Address Taichi (160 TOPS/W): argue why recurrent holographic beats feedforward chiplet for LLM token inference specifically

Requires lab (post-preprint)

- ☐ EXP-2, 3, 4, 5, 7 — experimental validation

Material Properties

ptr_glass

Photo-thermo-refractive glass for holographic recording

- **wavelength_range_nm**: [300, 3000]
- **max_refractive_index_change**: 0.005
- **absorption_coefficient_per_cm**: 0.01
- **diffraction_efficiency_max**: 0.99
- **write_wavelength_nm**: 325
- **erase_wavelength_nm**: 0
- **thermal_stability_C**: 400

| Cite: Glebov 2010, Proc. SPIE 7504, doi:10.1117/12.838767

linbo3_tfln

Thin-film lithium niobate electro-optic modulator platform

- **half_wave_voltage_V**: 1.5
- **bandwidth_GHz**: 100
- **propagation_loss_dB_per_cm**: 0.27
- **coupling_loss_dB**: 0.5
- **pockels_coefficient_r33_pm_per_V**: 30.9

| Cite: Wang et al. 2018, Nature, doi:10.1038/s41586-018-0551-y

sin_pic_a150

Ligentec A150 silicon nitride PIC platform, NIR optimized

- **wavelength_range_nm**: [700, 1060]
- **propagation_loss_dB_per_m**: 3.0
- **minimum_bend_radius_um**: 10
- **coupling_loss_fiber_to_chip_dB**: 1.5
- **platform**: Ligentec A150

| Cite: Ligentec A150 PDK documentation,
<https://www.ligentec.com/products/a150/>

gaas_vcsl_850nm

GaAs vertical-cavity surface-emitting laser at 850nm

- **threshold_current_mA**: 0.5
- **slope_efficiency_W_per_A**: 0.6
- **wall_plug_efficiency**: 0.35
- **linewidth_MHz**: 50
- **coherence_length_m**: 3.0
- **modulation_bandwidth_GHz**: 10

| Cite: Iga 2000, IEEE J. Sel. Top. Quantum Electron.,
doi:10.1109/2944.902166

ingaas_pin_detector

InGaAs PIN photodetector, telecom/NIR

- **responsivity_A_per_W**: 0.85
- **bandwidth_GHz**: 50
- **dark_current_nA**: 1.0
- **noise_equivalent_power_W_per_rthz**: 1e-14

| Cite: Bowers & Burrus 1987, J. Lightwave Technol.,
doi:10.1109/JLT.1987.1075507

si_pin_detector_850nm

Silicon PIN photodetector at 850nm

- **responsivity_A_per_W**: 0.6
- **bandwidth_GHz**: 10
- **dark_current_nA**: 0.1

silica_fiber_smf28

Corning SMF-28 single-mode fiber

- **attenuation_dB_per_km_1550nm**: 0.18
- **attenuation_dB_per_km_1310nm**: 0.35
- **attenuation_dB_per_km_850nm**: 2.5
- **core_diameter_um**: 8.2
- **numerical_aperture**: 0.14

| Cite: Corning SMF-28 Ultra datasheet, 2023,
| <https://www.corning.com/media/worldwide/coc/documents/Fiber/SMF-28%20Ultra.pdf>

gaas_vcse1_850nm_single_mode

Single-mode GaAs VCSEL at 850nm with narrow linewidth for coherent resonator operation

- **linewidth_MHz**: 10
- **coherence_length_m**: 30.0
- **threshold_current_mA**: 1.0
- **modulation_bandwidth_GHz**: 5

| Note: Oxide-confined single-mode VCSEL. Broader multimode VCSEL (50MHz, 6m l_c) insufficient for T>60 at L=20mm.

| Cite: Larsson 2011, IEEE J. Sel. Top. Quantum Electron. 17(6):1551, doi:10.1109/JSTQE.2011.2114837

Design Parameters

optical

Parameter	Value
wavelength_nm	850

token_embedding

Parameter	Value
dimension	512

resonator

Parameter	Value
geometry	Fabry-Perot
medium	PTR_glass
round_trips_T	100
cavity_length_L_mm	20
mirror_reflectivity_R	0.999

spatial

Parameter	Value
pixel_pitch_um	50
aperture_mm	2.5

snr

Parameter	Value
target_bits	6
target_snr_dB	38.0

interposer

Parameter	Value
detector	Si_PIN_850nm
amplifier	TIA_180nm_CMOS
nonlinearity	analog_comparator_GeLU
emitter	GaAs_VCSEL_850nm
latency_ns_per_layer	67
power_W_per_head	1.32

model

Parameter	Value
target_params	TBD
layers	TBD
embedding_dim	512

power

Parameter	Value
facility_W	TBD
target_per_token_mJ	TBD

Economic Analysis (5-Year TCO)

=====	
HYPERSCALE vs. QRI ECONOMIC COMPARISON	
=====	

CAPITAL COSTS ---	
Compute/Hardware	\$100.0B
\$800.0M	
Memory/Networking	\$42.5B
—	
Assembly	—
\$200.0M	
Control Electronics	—
\$2.0M	
Facilities	\$900.0M
\$38.5M	

```

-----
--
TOTAL CAPITAL                                $143.4B
$1.0B
Capital Advantage (ratio)      Hyperscale baseline      138x
cheaper

--- ANNUAL OPERATIONS ---
Energy                                $1.8B
$30.9M
Personnel                          $225.0M
$2.2M
Maintenance                        $28.7B
$40.0M
-----
--
ANNUAL OpEx                            $30.8B
$73.1M
OpEx Advantage (ratio)      Hyperscale baseline      422x
cheaper

--- 5-YEAR TOTAL COST OF OWNERSHIP ---
Capital                                $143.4B
$1.0B
OpEx ({} years)              $154.0B
$365.3M
-----
--
TOTAL 5-YEAR TCO                            $297.4B
$1.4B
TCO Advantage (ratio)      Hyperscale baseline      212x
cheaper

--- COST PER TOKEN (amortized over 5 years) ---
Cost per token                2.51e-05
1.19e-07
Token Cost Advantage      Hyperscale baseline      212x
cheaper

--- POWER & ENVIRONMENTAL IMPACT ---
Annual energy (kWh)          1.5e+10
3.5e+08
Annual CO2 (kg @ 0.4kg/kWh)      6.1e+09
1.4e+08
Carbon Reduction Ratio      Hyperscale baseline      44x
lower

--- THROUGHPUT EFFICIENCY ---
Total power draw (W)          1.8e+09
40000016.0
Tokens/sec per Watt          0.0
2
Efficiency Advantage      Hyperscale baseline      44x
better

--- PAYBACK PERIOD ---
Annual savings vs hyperscale      $30.7B
QRI capital investment          $1.0B
Payback period (100% efficiency)  0.4 months
Payback period (50% efficiency)  0.8 months
=====

```

Performance Scenarios

QRI PERFORMANCE: Baseline vs. Path A/B/C

Metric	Baseline	Path A (Cons)
Path B (Rec)	Path C (Agg)	
SNR (dB)	40.0	40.0
48.0	48.0	
NARG positions	1	16
128	128	
Throughput (M tok/s)	75.0	75.0
75.0	75.0	
Latency per-pos (ns)	13.30	0.831
0.104	0.104	
Latency gain (x)	—	16.0
128.0	128.0	
Parameters (M)	1.23	1.23
1.23	2.46	
Power (mW)	86.0	94.6
99.6	101.6	
Efficiency gain (x)	—	0.91
0.86	0.85	
Write overhead (%)	0	10
10	13	
Timeline (weeks)	0	2
6	8	
Cost (\$k)	0	0
5	6	
Ptychography	No	No
No	✓ Yes (2x)	

RECOMMENDATION: Path B (SNR upgrade + NARG 128)

- 8-16x latency improvement
- Realistic electronics (no exotic parts)
- Unlocks ptychography as well (marginal write SNR)
- \$5k cost, 4-6 week timeline
- Medium risk (TIA fab), medium-high reward

End of document. Generated by analyze/build_pdf_report.py on 2026-04-28.